

# Business intelligence

Unit 3 – Data exploitation. Query languages and visualization

S3-1 – OLAP

The seal of the University of Murcia is a circular emblem. It features a central figure, likely a saint or scholar, seated and holding a book. The figure is surrounded by three large, pointed Gothic arches. The entire scene is enclosed within a circular border containing Latin text. The text at the top reads 'UNIVERSITAS STUDIORVM MVRCIANAE' and the text at the bottom reads 'ANNO MCCLXXII'.

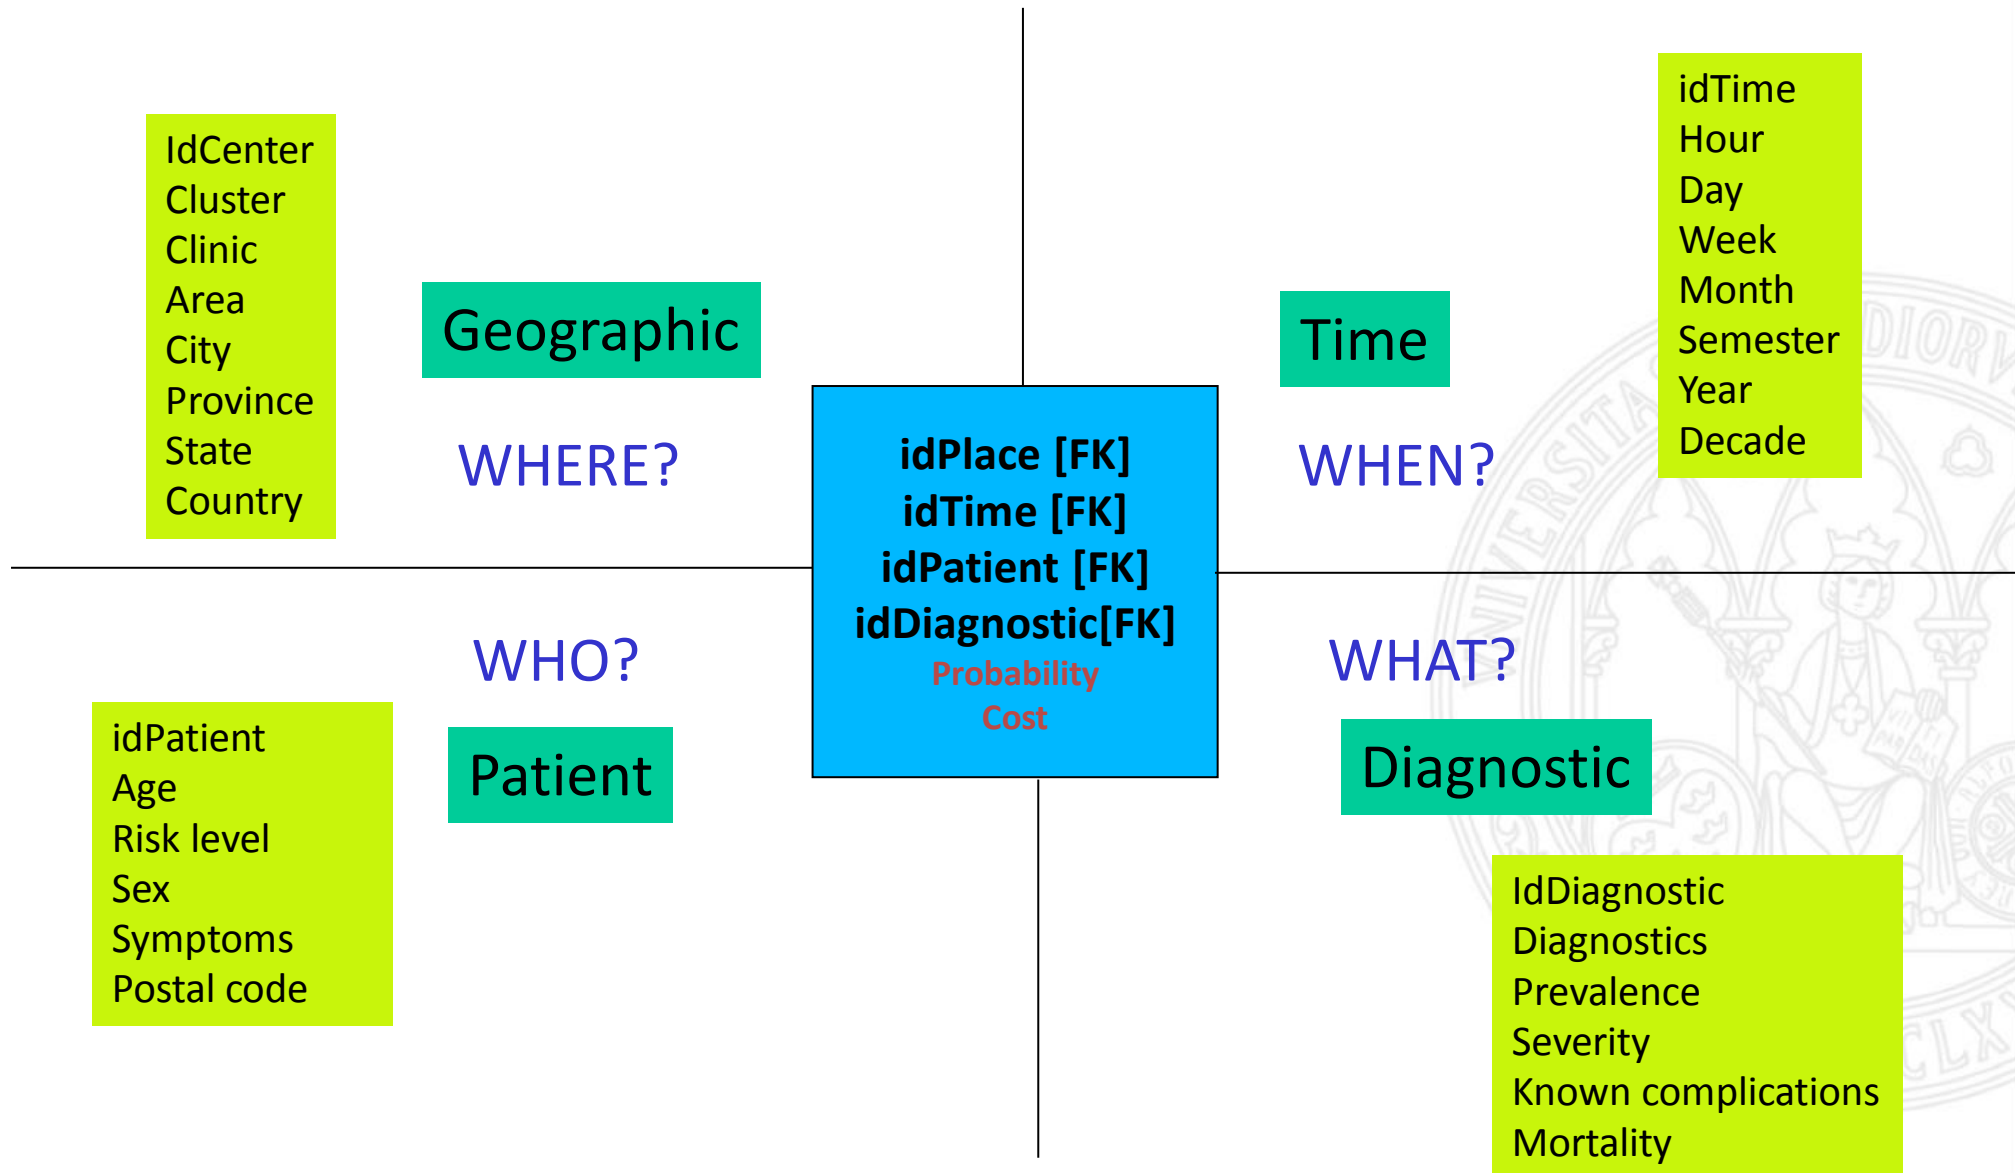
UNIVERSIDAD DE  
**MURCIA**

- OLAP tools provide the user with a multidimensional view of data (multidimensional schema) for each activity that is being analyzed.
- The user formulates queries to the OLAP tool selecting multidimensional attributes of this scheme without knowing the internal structure (physical schema) of the data warehouse.
- The tool generates a corresponding OLAP query and sends it to the query management system (eg by means a SQL SELECT statement).

- Query resolution procedure:
  - Build the query
  - Extract aggregated data
  - Visualize results
  - Analyze



- An OLAP query consists of
  - Retrieve **measures** or indicators
  - About the **facts**
  - **parametrized** by attributes in the dimensions
  - Constrained by **conditions** imposed on the dimensions
  - Eg: What is the **total cost** per **diagnostic** with **low mortality** rate **in the last year** for each **province** and **sex**?



Fact tables

Diagnostic	Sex	Total
D1	M	100
D1	F	200
D2	M	150
D2	F	75



2D view

Diag\Sex	M	F
D1	100	200
D2	150	75

Fact table

Diagnostic	Sex	Province	Total
D1	M	P1	100
D1	F	P1	200
D1	M	P1	100
D1	F	P1	200
D2	M	P2	150
D2	F	P2	75
D2	M	P2	150
D2	F	P2	75

3D view

Province

Diag\Sex	M	F
D1	100	200
D2	150	75

P2

Diag\Sex	M	F
D1	100	200
D2	150	75

P1

- The interesting thing is NOT ONLY to be able to query, in a way, something you can do with selections, projections, concatenation and traditional groupings.
- What is really interesting OLAP tools are its refinement operators for handling queries.
  - DRILL
  - ROLL
  - SLICE & DICE
  - PIVOT
  - ROLLUP
  - CUBE





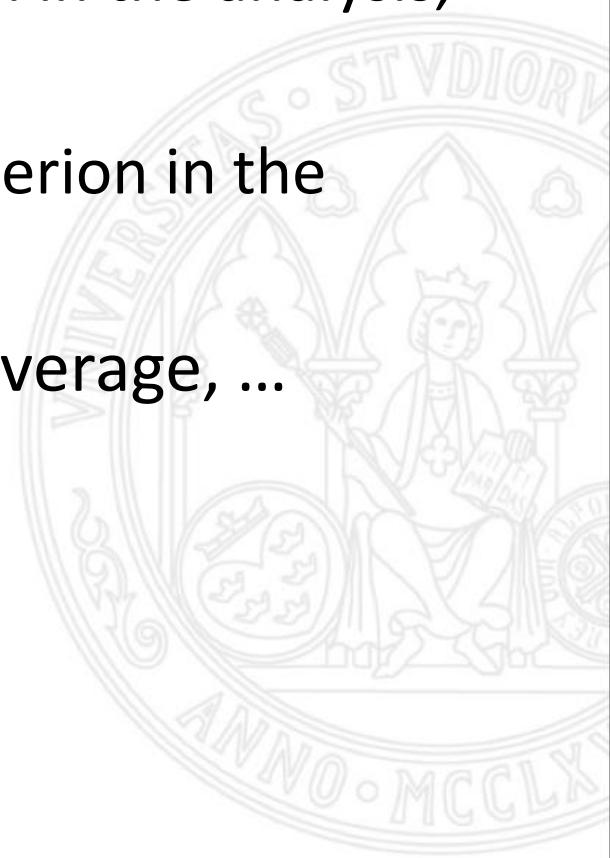
Diagnostic	Sex	Province	Total
D1	M	P1	100
D1	F	P1	200
D1	M	P1	100
D1	F	P1	200
D2	M	P2	150
D2	F	P2	75
D2	M	P2	150
D2	F	P2	75

Diagnostic	Sex	Total
D1	M	200
D1	F	400
D2	M	300
D2	F	150

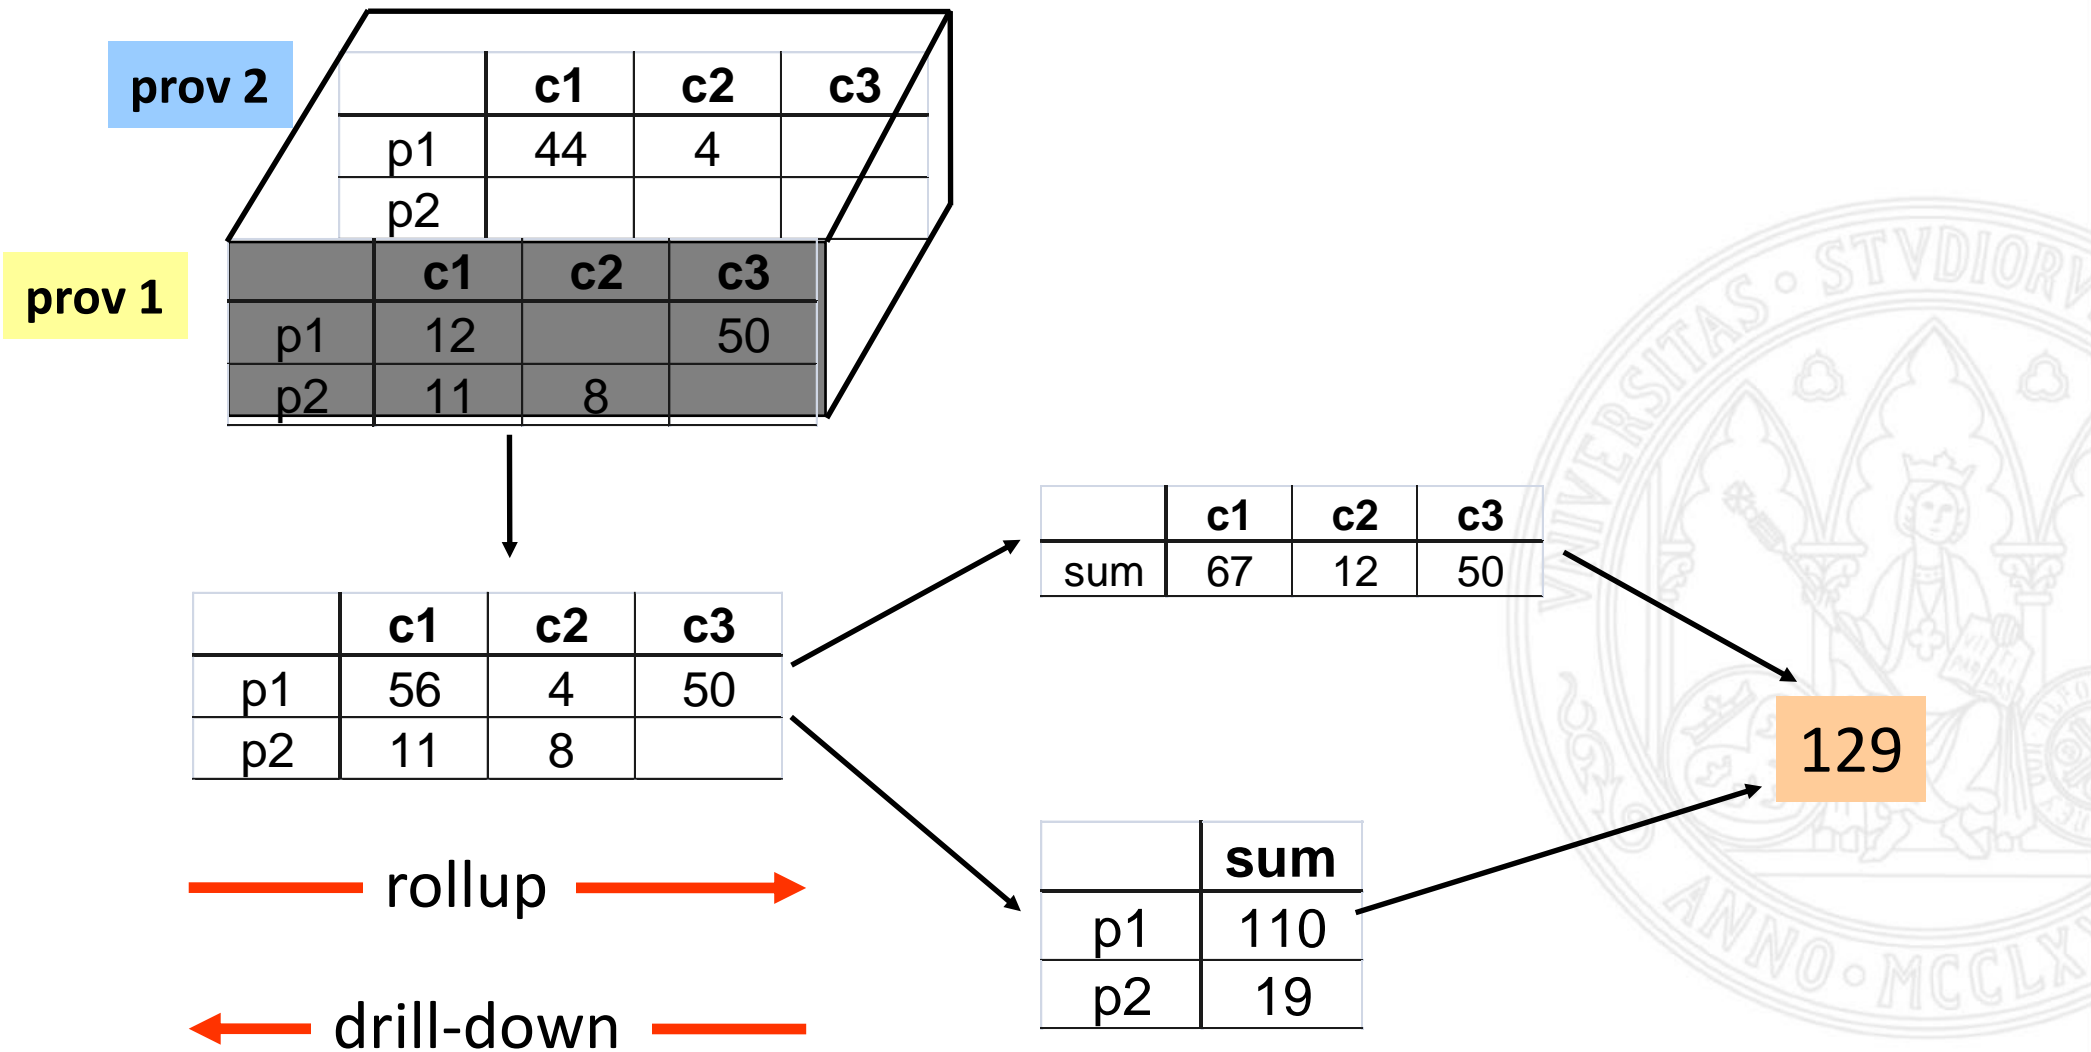
roll →

← drill

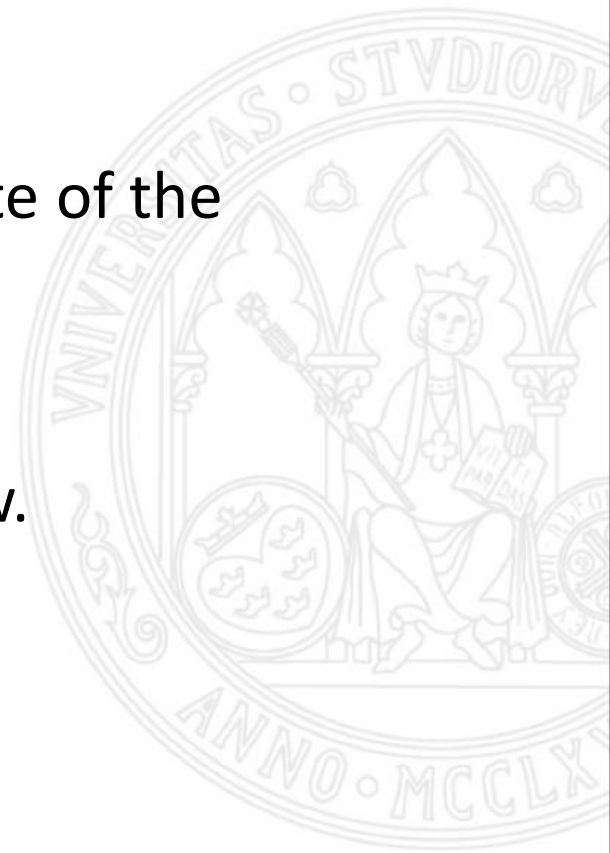
- Aggregate (consolidate) and disintegrate (division):
  - aggregation (**roll**): delete a grouping criterion in the analysis, aggregating the current groups.
  - disintegrate (**drill**): enter a new grouping criterion in the analysis, breaking existing groups.
- Aggregation in SQL: sum, count, max, min, average, ...



- DRILL (ROLL) can be done on:
  - attributes of one dimension on which a hierarchy has been defined:
    - **DRILL-DOWN**: upper to lower aggregation level
    - **ROLL-UP**: lower to upper aggregation level.
      - departament – category - product (Product)
      - year – semester – month – day (Time)
  - Other “drill”
    - **DRILL-ACROSS**: join several fact tables:
      - Careful implementation.
      - Sometimes this name is also used to change the dimension
    - **DRILL-THROUGH**: Use SQL to explore up to the relational back-end tables



- SLICE & DICE: select and project
  - SLICE: Delete a dimension on the analysis
  - DICE: Define a condition on some attribute of the dimension
- PIVOT: Rotate, reorientate the 2D- table view.



Diagnóstico	Sexo	Provincia	Total	Núm
D1	H	P1	100	6
D1	M	P1	200	5
D1	H	P2	100	6
D1	M	P2	200	11
D2	H	P1	150	7
D2	M	P1	75	7
D2	H	P2	150	2
D2	M	P2	75	1

Slice (Num) & dice (P1)

Diagnostic	Sex	Total
D1	M	100
D1	F	200
D2	M	150
D2	F	70


	Diagnóstico	Sexo	Total
P1	D1	H	100
	D1	M	200
	D2	H	150
	D2	M	75
P2	D1	H	100
	D1	M	200
	D2	H	150
	D2	M	75

Pivot



	Diagnóstico	Prov incia	Total
H	D1	P1	100
	D1	P2	100
	D2	P1	150
	D2	P2	150
M	D1	P1	200
	D1	P2	200
	D2	P1	75
	D2	P2	75

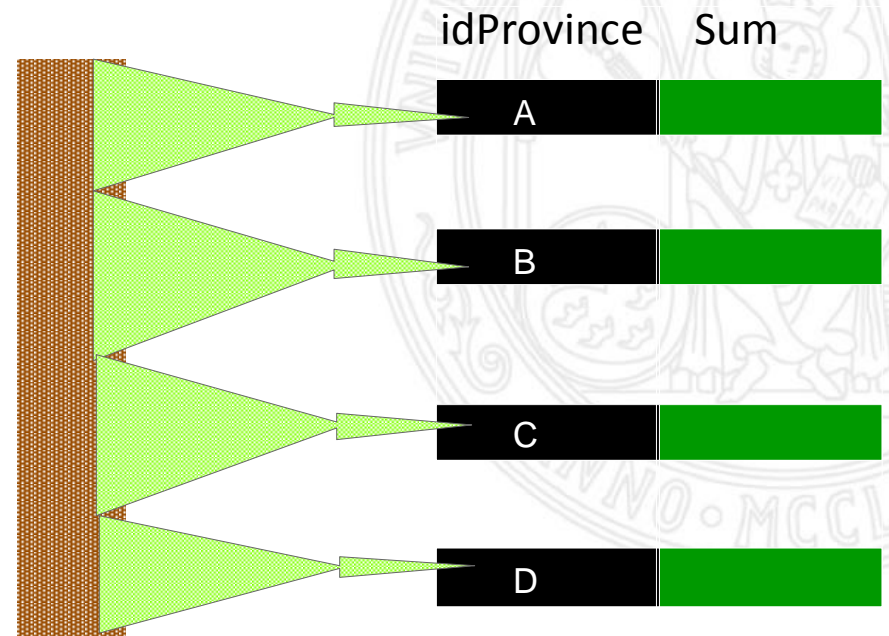
- SQL aggregation
  - `sum()`, `count()`, `avg()`, `min()`, `max()`
- Basic idea:
  - Combine values in one column
  - Into only one value
- Syntax:
  - `SELECT sum(cost) FROM diagnostic;`
- DISTINCT
  - Allows the aggregation only of different values
  - `SELECT COUNT(DISTINCT cost) FROM diagnostic`






- GROUP BY + HAVING
- Aggregating in subgroups of the table
- That fulfill some condition
- Syntax

```
SELECT idProvinc, sum(cost)  
FROM diagnostic  
GROUP BY idProvinc  
HAVING population > 2000;
```



- Limitations
  - Useful aggregations are difficult to calculate
    - Data cube
    - Complex: median, variance
    - Moving average
    - Rankings
  - Marginals or crosstabs
    - GROUP BY limited to 0-D and 1-D
  - Include sum and partial sums
    - drill-down & roll-up



- ROLLUP: performs the aggregation for the set of prefix of the attributes given
- Example:

```
SELECT item-name, color, size, SUM(number)
FROM sales
GROUP BY ROLLUP(item-name, color, size)
```

- Calculates SUM for the  $n+1$  prefixes:
- { (item-name, color, size), (item-name, color), (item-name), ( ) }
- Very useful for aggregating in hierarchies defined on dimensions
- It can be done in SQL without OLAP extensions, but very inefficiently.
- To improve efficiency: calculate the higher level aggregations using partial results of the more detailed levels

- CUBE: generalization of GROUP BY to n-dimensions.
- Calculates the aggregation function for all the subsets of the attributes given instead for only the prefixes (ROLLUP)
- Example:

```
SELECT item-name, color, size, SUM(number)
FROM sales
GROUP BY CUBE (item-name, color, size)
```

- Calculates the aggregate for the set of  $2^n$  combinations:
- $\{(item-name, color, size),$   
 $(item-name, color), (item-name, size), (color, size),$   
 $(item-name), (color), (size),$   
 $() \}$
- For each combination, the result is null for attributes that are not present in the combination.

- SQL:1999 uses NULL for representing both ALL and “usual” null
- In order to distinguish them we can use the GROUPING function that applied to an attribute
  - Returns 1 if NULL represents ALL
  - Returns 0 otherwise
  - Combined with DECODE we can return the desired value

- WINDOW clause defines **ordered** and **overlapping** groups of rows to calculate aggregates included at the end of each row.
- GROUP BY clause defines disjoint partitions of tuples in a sorted table, then calculates aggregates on those partitions, and generates a tuple with the result of the aggregate for each partition
- Example: "For each day, we want the average cost of obtaining diagnoses from the previous day, the current and the next, and cumulatively in the last 7 days":

SELECT date,

sum(cost) **OVER** (order by date **BETWEEN ROWS 1 preceding** and **1 following**),

sum(sum(precio) **OVER** (order by date **ROWS 7 preceding**))

FROM diagnostics;

- Syntax:
- SELECT attribute\_list\_1,
  - Aggregated\_function **OVER** W as windowName
- FROM table\_list
- WHERE constraints
- **WINDOW** W AS (
  - **PARTITION BY** attribute\_list\_2
  - **ORDER BY** attribute\_list\_3
  - frame declaration)



- Execution:
- FROM, WHERE, GROUP and HAVING generate an intermediate table.
- PARTITION: each partition contains tuples with the same values in the attributes given in attribute\_list\_2
- ORDER BY: rows in each partition are sorted according to the values of the attributes in attribute\_list\_3
- SELECT the tuples under the constraints established in the frame declaration
  - RANGE: logical conditions (ie: 5 days)
  - ROWS: in rows (ie: 5 preceding rows)



- Frame examples:
  - between rows unbounded preceding and current row
  - rows unbounded preceding
  - range between 10 preceding and current row
  - range interval 10 day preceding
  - range between interval 1 month preceding and interval 1 month following
- Default frame: If the frame is not specified, all preceding and current rows are considered in the partition
  - RANGE BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW

- RANK assigns to every tuple a rank based in some sorting of some attribute
- Example: given a cost-province relation rank each province by its cost.

```
SELECT province,  
       rank() over (order by coste desc) as provrank  
FROM diagnostic
```

- Afterwards, the result can be sorted by that field


```
SELECT province,  
       rank() over (order by coste desc) as provrank  
FROM diagnostic order by provrank
```

- RANKING allow gaps if there are 2 values with the same ranking.
  - Example: if the 1st and 2<sup>nd</sup> classified have the same cost, then both will be assigned rank 1, and the next row will have rank 3
  - DENSE\_RANK does not allow gaps, so the next row will have rank 2

- RANK over partitions:
  - “Rank the community and provinces by their cost”

```
SELECT province, comunidad,  
rank () over (partition by comunidad order by cost desc) as  
prov-comunidad-rank  
FROM diagnostic  
ORDER BY by comunidad, prov-comunidad-rank
```

- Several RANK can be included in the same query.



- Other rank functions
    - **percent\_rank**: it displays each row as a percentage of all the other rows up to 100% in a rank
    - **cume\_dist**: cummulative distribution
      - It displays the number of values in the set preceding and including x in the specified order divided by the number of rows.
    - **row\_number**
    - **ntile(x)**: cuantile
      - Divides the rows in the partition in x buckets with the same number ot rows
- ```
SELECT comunity, count (*),  
       ntile(3) over (order by count(*) desc) as quartile  
FROM diagnostic join patient  
GROUP BY comunity;
```

- Numeric functions (exp, cos, ln, ...)
- Aggregated (std, var, corr, regr, ...)
- Window functions:
  - Ranking: rank, dense\_rank
  - Distribution: percentage\_rank, cume\_dist
  - Count: row\_num
- Frame functions: lag, lead, ...
- SQL:1999 allows the use of **nulls first** and **nulls last**

```
SELECT student-id,  
rank ( ) over (order by marks desc nulls last) as s-rank  
FROM student-marks
```

- 1 Multidimensional view of data
- 2 Transparency to support (ROLAP, MOLAP)
- 6 Generic operations regarding the number of dimensions
- 9 Flexibility in the definition of the dimensions: constraints, aggregations and hierarchies among them.
- 10 Intuitive handling of operators: drill, roll, slice-&-dice, pivot.
- 12 No limit dimensions
- other:
  - 3 Accessibility from different data sources
  - 4 Coherent performance in reporting
  - 5 Client-Server Architecture
  - 7 Dynamic sparse matrix
  - 8 Multiuser support
  - 11 flexible report generation

